



RIVERDALE MIDDLE SCHOOL

UNESCO HACKATHON PROPOSAL · CRITICAL THINKING & AI EDUCATION / MEDIA & INFORMATION LITERACY

I Thought I Knew You

A critical-thinking and AI-literacy game disguised as a four-night group-chat mystery — teenagers learn to reason about a video cut to mislead, a misdated photo, a cloned social account, and a cloned voice, using a friendship they care about as the reason to bother.

TEAM

— add your names here —

TRACK

— add hackathon track —

FORMAT

Browser game · client-side AI

AUDIENCE

Secondary school, ages 13+

What's in this document

A game about one thing: teaching teenagers to reason carefully about AI-manipulated media, using a friendship they care about as the reason to bother.

01	The Problem	04
02	Theoretical Background the MIL Index model	05
03	Theoretical Background why people get fooled	06
04	Theoretical Background why it's a story	07
05	Our Solution	08
06	How It Works	09
07	How It Works day by day	10
08	The Notebook	11
09	One Mechanic, Close Up	12
10	Technology	14
11	Educational Impact	15
12	Alignment with the MIL Index	16
13	Roadmap	17
14	Team & Contact	18

What this project sets out to do

A critical-thinking and AI-literacy game that gives teenagers one habit — checking AI-manipulated media before they react — practised through a story they care about, not taught through a quiz.

The problem

Cyberbullying among school students is widespread and rising — and generative AI has made it worse. Manipulated "evidence" (a cut video, a cloned voice, a fake account) now spreads through a class group chat faster than anyone stops to check.

The objective

Give students one habit: pause to check AI-manipulated media before they react — turning "check before you forward" from a rule they're told into a reflex they've rehearsed.

The approach

A four-night group-chat mystery. Real, checkable manipulations; real social consequences; a friendship as the reason to bother — practised, not lectured.

Who it's for

Secondary-school students, ages 13+ — met where manipulated media actually reaches them: forwarded into a class group chat by someone they trust. Trains all five pillars of media & information literacy.

Success isn't a big score — it's a student who instinctively pauses before forwarding the next thing that lands in their group chat.

You don't need real evidence to hurt someone anymore

Cyberbullying among school students is already widespread — around 1 in 6 teenagers have experienced it in school, and the figure keeps rising, drawn from more than 279,000 adolescents across 44 countries (WHO/Europe, 2024). And unlike bullying inside the school gates, it doesn't end when a student gets home; it follows them onto the phone in their bedroom. It's also rarely a stranger doing it — it's usually a classmate, someone inside the same circle of trust, and much of what they spread isn't even true. More than one in five teenagers say false rumours have been started about them (Pew Research Center, 2022): a story that never happened, pushed around the class until it sticks.

Generative AI has made that far easier. Now that any student can fake a voice, a photo, or a video in seconds, roughly one in five secondary schools have already seen a fabricated clip or image turned against a student (RAND, 2024), and two-thirds of school staff say their students have been fooled by one (EdWeek Research Center, 2024). And not all of it is even AI — sometimes it's real footage stripped of its context. Either way, it lands in the class group chat, about a classmate everyone knows, with the group reacting before anyone has checked.

STRIPPED
CONTEXT

**Real
media,
false
frame**

A genuine clip or photo, cut or re-dated so it means something it didn't mean.

SPOOFED
IDENTITY

**Fake
accounts,
cloned
voices**

AI-generated images and synthetic audio built to pass as someone real.

SOCIAL
PRESSURE

**Reacting
before
checking**

Group chats reward speed of reaction, not accuracy — verification feels slow and socially costly.

None of these three failure modes are solved by teaching young people to "spot the fake" in isolation — the video in this scenario, for instance, isn't fake at all. What's missing isn't detection skill in the abstract; it's the **habit** of pausing, tracing a claim back to its source, and asking who benefits from you reacting fast — the same four-move discipline (Stop, Investigate the source, Find better coverage, Trace claims to the original context) known as the **SIFT method**, and this project's entire investigation mechanic was built around it (Caulfield, 2019). That is a behavioural gap, not a knowledge gap — which is why this project is built as a game about consequences, not a quiz about facts.

Real techniques, modelled directly

Every piece of "evidence" in the game is built from a manipulation technique already circulating online today — nothing in the scenario is speculative or futuristic.

Selective re-cutting of real video

Recycled photo, new date/caption

AI face-swapped impostor account

Cloned voice message

False "AI-authentic" confidence badge



The framework: DW Akademie's MIL Index

The project's primary goal is critical thinking and AI literacy — the friendship story is the delivery mechanism, not the message. That goal is built on a named competency framework for what "critical thinking about media" actually consists of: this hackathon's own reference model for media and information literacy.

1 • The MIL Index model



- **Access** — using media technology to reach a message in the first place, and knowing where to look for information.
- **Analyze** — interpreting and critically evaluating a media message against what you already know.
- **Reflect** — critically examining which sources were actually used, and what impact that choice had.
- **Create** — composing and sharing a message of your own to express an idea or opinion.
- **Act** — putting the skill into practice, for your own benefit and your community's.

Diagram and definitions adapted from DW Akademie's Media and Information Literacy Index, the "MIL Heroes" model (DW Akademie, 2020) — the reference framework behind this hackathon's "Heroes and Villains in the Age of AI" brief. Full pillar-by-pillar mapping to specific game mechanics appears in section 12, "Alignment with the MIL Index." Full citation on page 6.

Most "media literacy" games only exercise Analyze — spot the fake. This one was designed to exercise all five pillars in a single playthrough, not detection in isolation.

Why this exact model, and not a generic one


The MIL Index fits this project because its five pillars mirror the exact sequence a player lives through when a manipulated video lands in a group chat: they first have to *access* it before they can react, *analyze* what looks off, *reflect* on who sent it and why, *create* a response, and *act* on what they decide. A framework built around one skill — spotting the fake — would only cover the second step; this one gives the game a full arc instead of a single detection puzzle. It also keeps the design honest: treating "critical thinking" as five distinct, separately-trainable competencies, rather than one vague instinct, is what makes the game's design falsifiable — each pillar below is tied to a specific, checkable mechanic (section 12), not a vibe.

Why people get fooled, and how practice fixes it

Knowing the MIL pillars explains *what* to train. These two ideas, both backed by research in misinformation studies, explain *why people get fooled* — and *why rehearsal is what fixes it*.


2 · The psychology of misinformation

People don't fall for false claims because they lack judgment — they fall for them because memory works in a specific, predictable way. This is called **source monitoring**: the ability to remember *where* a belief came from, not just what the belief is. People hold onto the content of a claim far longer than they hold onto its source, so once the source is forgotten, a plausible-sounding claim starts to "feel true" on its own (Johnson & Raye, 1981). The same mechanism explains the "misinformation effect," where a false detail gets attached to a real memory as if it always belonged there (Lindsay & Johnson, 1984). None of this is about being careless or clever — it makes critical thinking a specific, learnable habit: reconnecting a claim to its source before reacting to it. That habit is exactly what the game's investigation actions train.




3 · Prebunking / inoculation theory

The idea behind inoculation theory is the same one behind a medical vaccine: give someone a weak, safe dose of a manipulation technique, then immediately explain the trick — and their resistance to the real thing gets stronger. Controlled studies of this approach, including the browser game *Bad News*, found that people pre-exposed to weakened misinformation techniques got measurably better at spotting manipulation afterwards, and the effect held up across different cultures (Roozenbeek & van der Linden, 2019; Basol et al., 2021). Each of the five scenarios in the game follows this same expose-then-explain structure, repeated until it becomes a reflex instead of a rule the player has to remember:



Both ideas, in one game moment



Why it's a story, and why it's a game

The MIL pillars and the psychology of belief explain what to train and why rehearsal works. These two ideas explain why that rehearsal is wrapped in a first-person story instead of a lecture or a checklist.

4 · Narrative transportation & perspective-taking

Audiences who become absorbed, or "transported," into a story pay closer attention, form vivid mental imagery, and are measurably more likely to adopt the beliefs and attitudes the story implies than audiences given the same information directly (Green & Brock, 2000). This is why the vehicle is a friendship, not a worksheet. Wednesday night pushes this further and deliberately reverses the camera — the player becomes the one being falsely accused — turning "I can judge Nicole's evidence" into "I understand exactly how this gets done to someone," a stronger, stickier form of the same critical-thinking lesson.

5 · Experiential / game-based learning

Kolb's experiential learning cycle holds that durable learning comes from a loop of concrete experience, reflection, and active experimentation, not from being told a conclusion (Kolb, 1984). Games are a natural fit for that loop: Gee's study of learning in well-designed video games found that they routinely embed exactly this structure — act, get real feedback, adjust — better than most classroom formats manage (Gee, 2007). The game never states a rule and asks the player to recall it; it withholds the answer, lets a real (in-fiction) consequence follow from the player's actual choice, and only then shows what was true.

Teaching "how to spot a deepfake" is a knowledge problem. Teaching someone to feel the urge to check before they forward something is a behavioural one — and behaviour is what rehearsed, consequence-driven game mechanics are good at shaping.

IDEA	WHERE IT ACTUALLY SHOWS UP IN THE GAME
Narrative transportation	The whole friendship framing, and Wednesday night's camera reversal onto the player.
Experiential learning	No stated rules — only actions, consequences, and a debrief that shows what was true.

References

Caulfield, M. (2019). *SIFT (The Four Moves)*. Hapgood. hapgood.us/2019/06/19/sift-the-four-moves

DW Akademie. (2020). *Media and Information Literacy Index — "MIL Heroes."* Deutsche Welle Akademie.

Johnson, M. K., & Raye, C. L. (1981). Reality monitoring. *Psychological Review*, 88(1), 67–85.

Lindsay, D. S., & Johnson, M. K. (1989). The eyewitness suggestibility effect and memory for source. *Memory & Cognition*, 17(4), 349–358.

Roozenbeek, J., & van der Linden, S. (2019). Fake news game confers psychological resistance against online misinformation. *Palgrave Communications*, 5, 65.

Basol, M., Roozenbeek, J., Berriche, M., Uenal, F., McClanahan, W. P., & van der Linden, S. (2021). Towards psychological herd immunity: Cross-cultural evidence for two prebunking interventions against COVID-19 misinformation. *Big Data & Society*, 8(1).

Green, M. C., & Brock, T. C. (2000). The role of transportation in the persuasiveness of public narratives. *Journal of Personality and Social Psychology*, 79(5), 701–721.

Kolb, D. A. (1984). *Experiential Learning: Experience as the Source of Learning and Development*. Prentice-Hall.

Gee, J. P. (2007). *What Video Games Have to Teach Us About Learning and Literacy* (2nd ed.). Palgrave Macmillan.

EdWeek Research Center. (2024). *Educator and student survey on deepfakes in schools*. Education Week.


Pew Research Center. (2022). *Teens and Cyberbullying 2022*. Pew Research Center. [pewresearch.org/internet/2022/12/15/teens-and-cyberbullying-2022](https://www.pewresearch.org/internet/2022/12/15/teens-and-cyberbullying-2022)

RAND Corporation. (2024). *Artificially Intelligent Bullies: Dealing with Deepfakes in K–12 Schools*. RAND. [rand.org/pubs/research_reports/RRA3930-5.html](https://www.rand.org/pubs/research_reports/RRA3930-5.html)

World Health Organization Regional Office for Europe. (2024). *A focus on adolescent peer violence and bullying in Europe, central Asia and Canada* (HBSC 2021/2022 survey, Vol. 2). WHO/Europe.


Four nights, five reasoning exercises, one friendship

Structurally, this is five separate exercises in evaluating AI-manipulated media, each with a different manipulation technique and a different detectable tell. Narratively, it's one continuous story across four evenings: the player is one of 28 classmates in a group chat, and on Monday, Tuesday, and Wednesday the chat receives one piece of damning "evidence" against Nicole, the player's best friend of seven years — each staged as something that just landed on your phone, not as a lesson. Wednesday night, before the story even reaches Thursday, the same trick turns on the player themselves.



Every "tell" is real and checkable — a full-length video that recontextualises the clip, a dated post that recontextualises the photo, a visual or stylistic artefact in the fake account, a room-tone difference in the voice message. Nothing requires the player to be told the answer; it requires them to look. The title pays off literally at the end of night one, when a classmate posts: *"wow. i thought i knew her."*

Why a friendship, not a headline



Every one of these five scenarios resolves the same way: the truth was checkable the whole time, and the player either checked or didn't. Nothing in the story is unknowable — that's a deliberate design constraint, not a coincidence.

Investigate. React. Or say nothing at all.

Each evening is a self-contained reasoning task dressed as a phone screen. The player has limited time to spend on three kinds of moves: **investigate** (reverse-image search, ask a witness, compare accounts, replay audio slowed down) — the actual critical-thinking work — **react in the group** (forward, pile on, defend, stay silent), or **message Nicole privately**. Every choice moves a friendship score, the group's opinion of the player, and which of several endings is reached, so reasoning well and reasoning fast are put in direct, felt tension.

NIGHT	EVIDENCE	WHAT'S ACTUALLY TRUE	THE DETECTABLE TELL
Monday	A party video	Real footage, real party — but cropped before the sentence that explains it.	Ask for the full clip.
Tuesday	A party photo	Completely real and unedited — just from a party a year earlier.	Check the original post date.
Wednesday	A second account	Not her — an AI-generated impersonation. Her real, silent account is still live.	Compare the two accounts.
Wed, that evening	A voice clip of you	An AI clone of the player's own voice, spliced from a real recording made days earlier.	Compare it to the original.
Thursday	A voice message	An AI voice clone; the "97% authentic" checker badge is simply wrong.	Listen for room tone twice.

The chat, in character

Every classmate holds a fixed emotional stance, delivered live through a small AI language model running in the player's own browser (see *Technology*, next). No two conversations play identically, but no character ever contradicts the true story.



Nicole

The accused. Direct, casual, imperfect — worried, but not performing it.



Hanna

Certain, dismissive, quick to pile on. Opposes Nicole from night one.



Lea

Follows Hanna's read of the room, rarely first to speak.



Mia

The skeptic. Questions the evidence if the player gives her a reason to.



Benito

Evasive. Says little — but he knows more than he says.

The three moves, and what each one trains

ACCESS + ANALYZE Investigate Costs real time — the reasoning work.	SOCIAL PRESSURE React in the group Instant, public, and free — tempting.	CREATE + ACT Message Nicole Private and direct, away from pressure.
---	---	--

Every night's actual reasoning task

Each investigation action below is a real, selectable move from the game — the exact label the player sees, and what checking it actually reveals. Some are genuine progress; a few are deliberate dead ends, because knowing when verification effort isn't paying off is part of the skill.

 Monday — the video DISGUST · REAL FOOTAGE, CUT BEFORE THE START

- **"Watch it again, slowly"** — the clip starts mid-word, proving something was cut.
- **"Ask Mia, she was at that party"** — she confirms Nicole was quoting someone else.
- **"Ask who filmed it"** — Hanna admits she doesn't know, just received it.
- **"Reverse search the video"** — dead end, no matches (costs time, teaches nothing).

 Tuesday — the photo ANGER · GENUINELY REAL, JUST FROM A YEAR EARLIER

- **"Reverse search the photo"** — the same file was uploaded back in July.
- **"Look at the background"** — a half-hidden banner reads "LAURA'S BIRTHDAY."
- **"Look at the photo properly"** — bare arms outdoors at night; wrong for November.
- **"Ask Nicole about the photo"** — she hasn't left her room in two days: an alibi.

 Wednesday — the second account VIOLATION · AI IMPERSONATION

- **"Put the two accounts side by side"** — real: 52 posts since 2022; fake: 4 posts, all this week.
- **"Compare the writing to her real captions"** — real Nicole never uses full stops or capitals; the fake does both.
- **"Look at the followers"** — the fake's 42 followers all followed this week.
- **"Ask Nicole if the account is hers"** — she denies it directly.

 Wednesday, that evening — the clip of you RECOGNITION · YOUR OWN WORDS, SPLICED

- **"Compare it to the original"** — matches your real Sunday voice note; same words, different order.
- **"Listen again, slowly"** — a phrase cuts off mid-word instead of trailing off naturally.
- **"Ask when it was recorded"** — Nicole herself notices the timestamp doesn't add up.

 Thursday — the voice message CERTAINTY · CLONED VOICE, "97% AUTHENTIC" CHECKER IS WRONG

- **"Listen closely, twice"** — no room tone, no breathing, identically-timed pauses.
- **"Ask where she was tonight"** — Mia confirms Nicole was with family all night.
- **"Compare it to her old voice notes"** — real ones trail off and laugh; this one never does.
- **"Run the AI checker yourself"** — returns "97% AUTHENTIC" — a deliberate false-confidence trap.

Time is the real currency


Every action above costs limited in-game time, so "check everything" is never free — reasoning well under a time constraint is the actual skill being trained, not reasoning well with unlimited time. The debrief on the final night (see *Educational Impact*) measures exactly where that time went.

- Investigate — costs time
- React in the group — instant, public
- Message Nicole — private, direct
- Report an account — reversible consequence

A desk, not a debrief screen


Reflection isn't a quiz that appears once, at the end of the week. It's a physical desk the player can open any evening, holding three books that fill in as the investigation actually progresses.

Diary




One tab per night, each holding the player's own verdict note on what they believed and why — written by the player, not chosen from options.

Media Literacy Handbook



The same five MIL pillars from section 2, restated in-fiction as the player's own field notes rather than a framework slide.

Class 10b Album



A relationship map of the class that starts mostly blank — nodes and edges unlock only as the player actually learns who stands where.


Because the map and handbook only fill in with what's actually been investigated, checking the "Access" and "Analyze" pillars (see section 2) leaves a visible, growing trace the player can revisit — reflection becomes something they build, not something the game tells them happened.

Wednesday: which account is real?

This is the actual in-game comparison screen from the "second account" scenario — the clearest single example of what the player is asked to practise: not "trust your gut," but look for the specific, findable tell.

Nicole's real account


4 YEARS OLD



- Long, ordinary posting history
- Photos and captions match her established style
- Has gone quiet — she isn't posting to defend herself

The copycat account

DAYS OLD



- Same name and photos, wrong writing voice
- AI-image artefacts on close inspection
- Posting fast, right when it's most damaging

The two screens are visually similar on purpose — this is the point. Judges and players alike are meant to have to actually compare them side by side, the same way the in-game "compare accounts" action works, rather than intuit which is fake at a glance.

What testing changed

Before calling the game finished, we put the prototype in front of players from the target group — and rebuilt it around what we saw.

Photo: playtesting the prototype with a target-group player — placeholder

How we tested

We watched three players from the target group work through the game, noting both what landed and where they got stuck. Those pain points shaped what we changed next.

What we changed

- Players felt lost at first, unsure what to do — so we added the notebook and clickable hints to give them a clear sense of where to look next.
- The prototype never showed the full story, so players couldn't tell what was actually true. Now, once they reach 100%, the game reveals the whole situation.
- Players fixated on finding the culprit, not on which evidence was real or fake — so each day now ends with a "real / real-but-cut / fake" checklist.

What worked

Players were positively surprised by the story and liked discovering how to check information for themselves. Day 3 — their own voice cloned against them — caught them off guard.

"I didn't expect the game to use my own voice against me."

— Playtester 1

A limitation we're addressing

Our ongoing challenge is keeping players focused on analysing the evidence rather than on *who did it* — the per-day checklist is a first step we're still sharpening.


An AI that runs on the player's device, not a server

The in-character chat is powered by a small, quantised language model (LFM2.5-350M, GGUF Q4_K_M, ~219 MB) run entirely client-side in the browser via WebAssembly, with optional WebGPU acceleration. Every reply is grounded against a short, per-day fact sheet the player never sees, so improvised dialogue can never contradict the "true" version of events. This is also, quietly, AI education in itself: the player is talking to a real, working generative model the whole time — a concrete, hands-on referent for "what an AI actually is" underneath every abstract warning about deepfakes.

Why that matters for schools

- **No account, no server, no data collection** from minors — nothing typed by a student leaves their own device.
- **Works after one load**, even on slow or filtered school networks, since there is no ongoing API call per message.
- **No build step.** Plain static files — deployable from a USB stick, a school server, or free static hosting, with zero infrastructure cost.
- **Content lives in editable JSON**, separate from code — a teacher, translator, or student team can extend or localise scenarios without touching a line of logic.

Typical AI chat vs. this game



Model card

Model	LFM2.5-350M (LiquidAI), GGUF Q4_K_M quantisation
Footprint	~219 MB, downloaded once and cached in the browser
Runtime	wllama (llama.cpp compiled to WebAssembly), optional WebGPU, CPU fallback
Grounding	A short per-day fact sheet constrains every reply so improvisation can't contradict the true story
Guardrails	Replies breaking character or sounding like an assistant are rejected and retried
Voice clone	A second on-device model, Pocket TTS (Kyutai, ~146 MB ONNX, in a Web Worker) — clones the player's own voice for Wednesday night's "clip of you"

The lesson is the ending, not a warning label

There is no "did you spot the deepfake?" quiz. Instead, the game silently tracks what the player actually did all week, and hands it back to them, unfiltered, at the end.

What gets measured

A real excerpt from the in-game end-of-week ledger:

Forwarded	— things
Reacted	— times
Reported her real account	yes / no
Reported the fake account	yes / no
Time spent checking vs. reacting	— min / — min

What gets said back to you

Real debrief lines, selected by what actually happened in the playthrough:

"You never looked for the original of anything. Originals existed."

"You never compared the account to her real one. It was still up."

"You checked all five. That's rarer than it should be."

The per-check debrief lines above are composed from exactly what was and wasn't verified, so the specific feedback differs player to player. Which of the game's six authored endings closes the week is then decided by a signal ledger — checks completed, public support, private messages to Nicole, and the resulting relationship — not a simple multiple-choice branch. Endings the player has earned are kept in an in-game gallery, so a completionist can see how differently the week could have gone.

Six endings, chosen by what you did

WORST

I thought I knew you

"I thought I knew you, {name}."

LET THEM KEEP TELLING IT

"You found the crack in their story. You let them keep telling it."

WORDS IN THE DARK

"A few words in the dark. None in the light."

KEPT IT PRIVATE

"You were the only one who knew. And you kept it between the two of you."

TOLD THE ROOM


"You told the truth to the room. The room grew quieter around you."

SPOKE IT ALOUD

"You waited for the truth. Then you spoke it aloud. She heard you twice."

Every pillar, one concrete mechanic each

Section 2 introduced DW Akademie's MIL Index — the five-pillar model this hackathon's own brief is built around. Rather than claim generic alignment, here is exactly which game mechanic exercises each pillar, in the order a player actually encounters them: detect first, reflect after, then create and act.



In one sentence: this is a MIL Index exercise disguised as a story about a friendship — every pillar exercised, in the same order, every single playthrough.

Equity note: none of the above requires a server, an account, or reliable connectivity after first load — which matters for Access specifically, since a pillar about "knowing where to find information" is undermined if the tool itself is hard to reach in an under-resourced classroom. This describes thematic alignment with DW Akademie's published MIL Index model; it is not a claim of formal review, partnership, or endorsement.

From hackathon build to classroom tool

1 Facilitator guide

A short companion document for teachers: discussion prompts per night, and how to read the anonymised end-of-week ledger as a class discussion starter rather than a grade.

2 Localisation

All narrative content already lives in separate JSON files, decoupled from game logic — translation into other languages is a content task, not an engineering one.

3 Accessibility pass

Screen-reader support for the chat interface, captions for the voice-message evidence, and colour-contrast review of the investigation UI.

4 More scenarios

Additional evidence types (manipulated group photos, forwarded-chain misinformation) reusing the same investigate / react / debrief structure.

5 Aggregate, opt-in classroom view

An anonymised summary of common omissions across a class set, to focus a follow-up lesson on the habits a specific group actually struggled with.

Static hosting · GitHub Pages


No build step

Content in JSON

Client-side LLM

No student data collected

Where this stands



Who built this

NAME	ROLE	CONTACT
Mine	Design	— add contact —
Adrian	Programming	— add contact —
Carlos	Programming	— add contact —
Karin	Design	— add contact —

Team photo, logo, or additional credits — placeholder

Try it / find it

Live game

barriletechapin.github.io/I-thought-I-knew-you

Source

github.com/BarrileteChapin/I-thought-I-knew-you

"I thought I knew you."

— the line every ending in the game earns differently, depending on what the player actually checked.

Thank you for reading. Questions welcome before, during, or after the pitch.